TRAINING LLMs

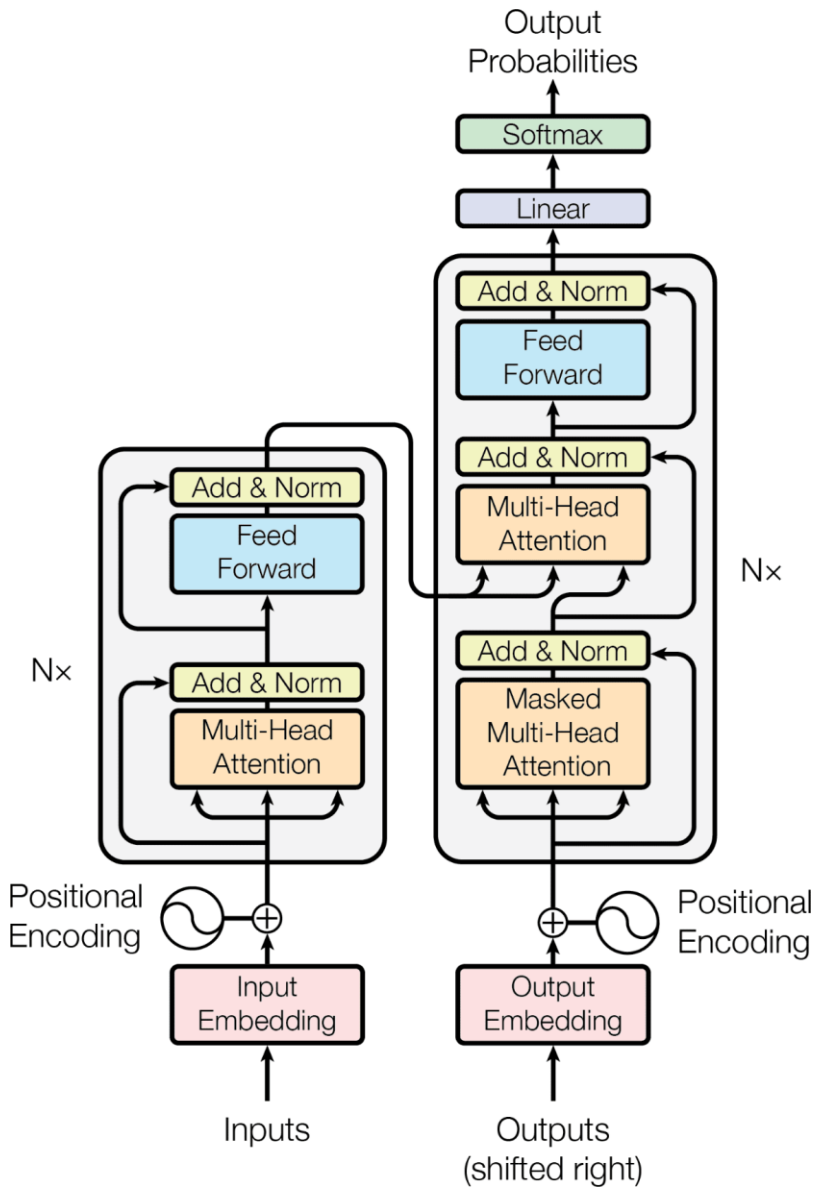

Transformers

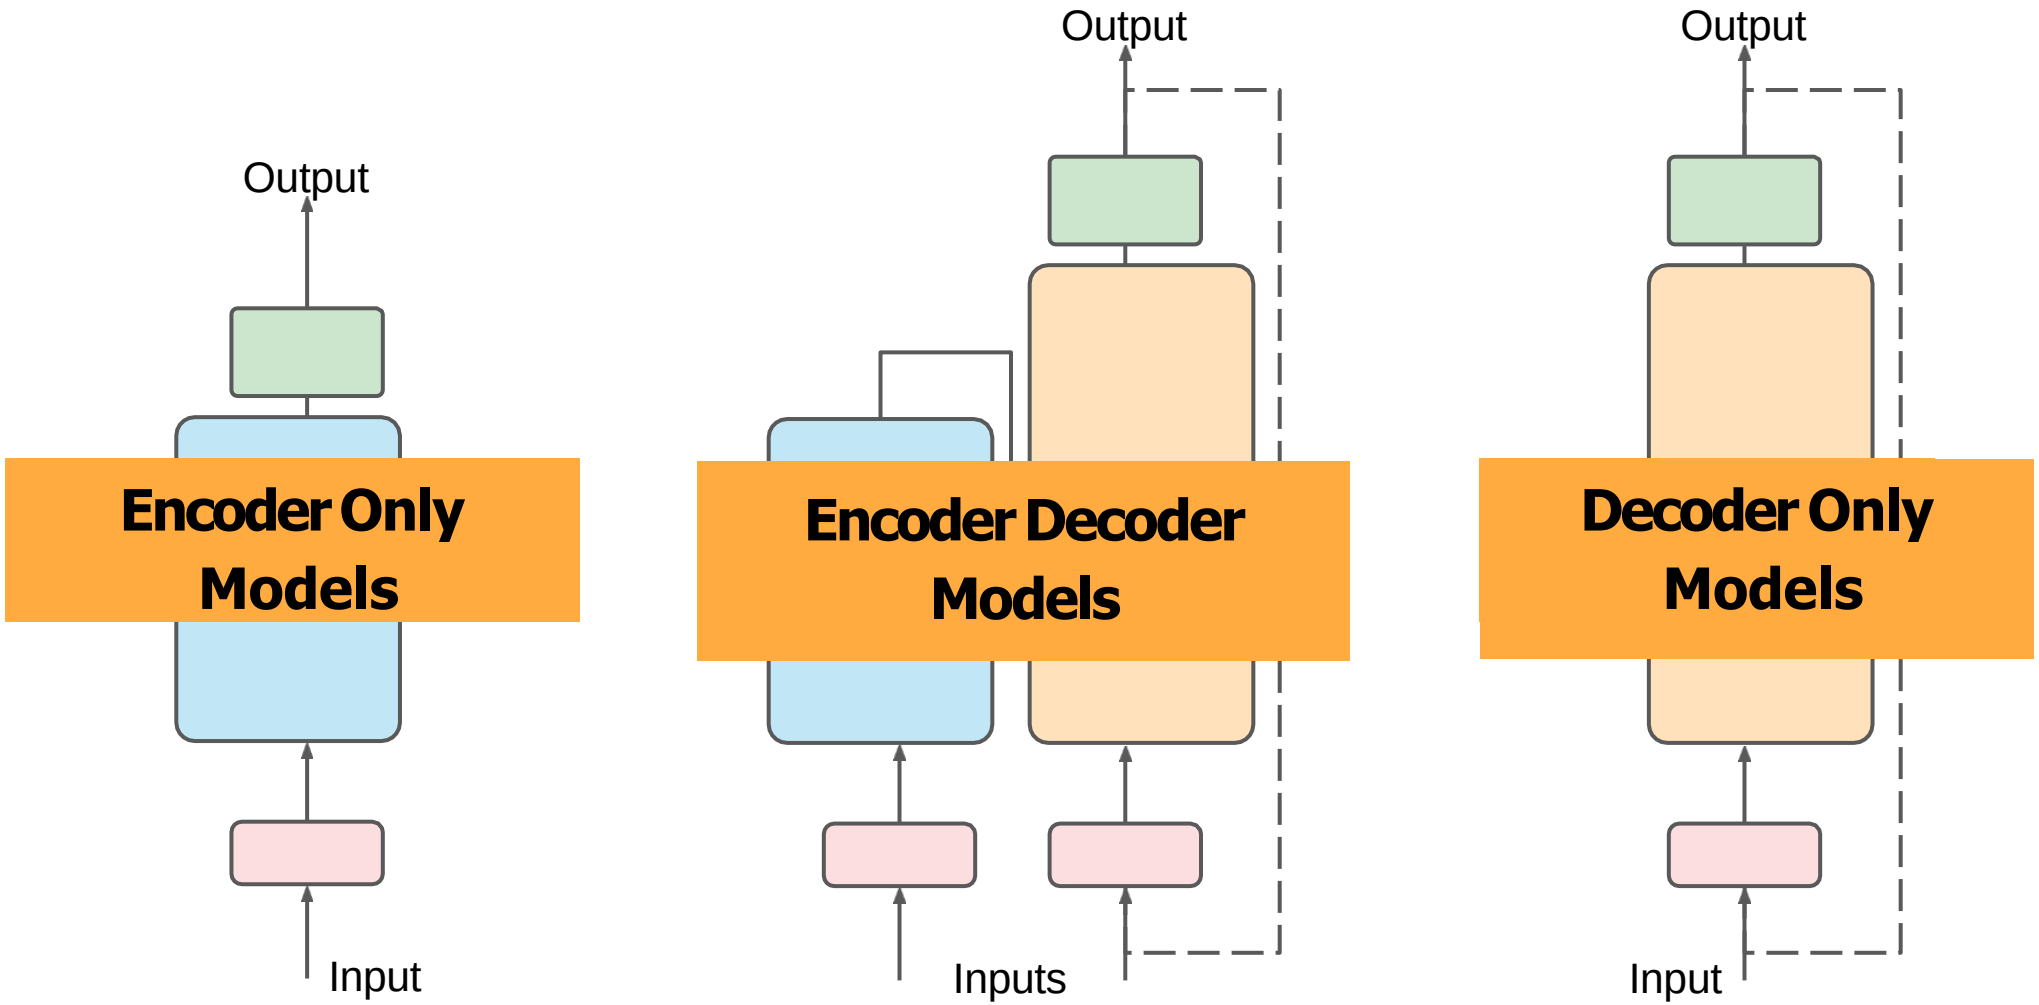

Transformers

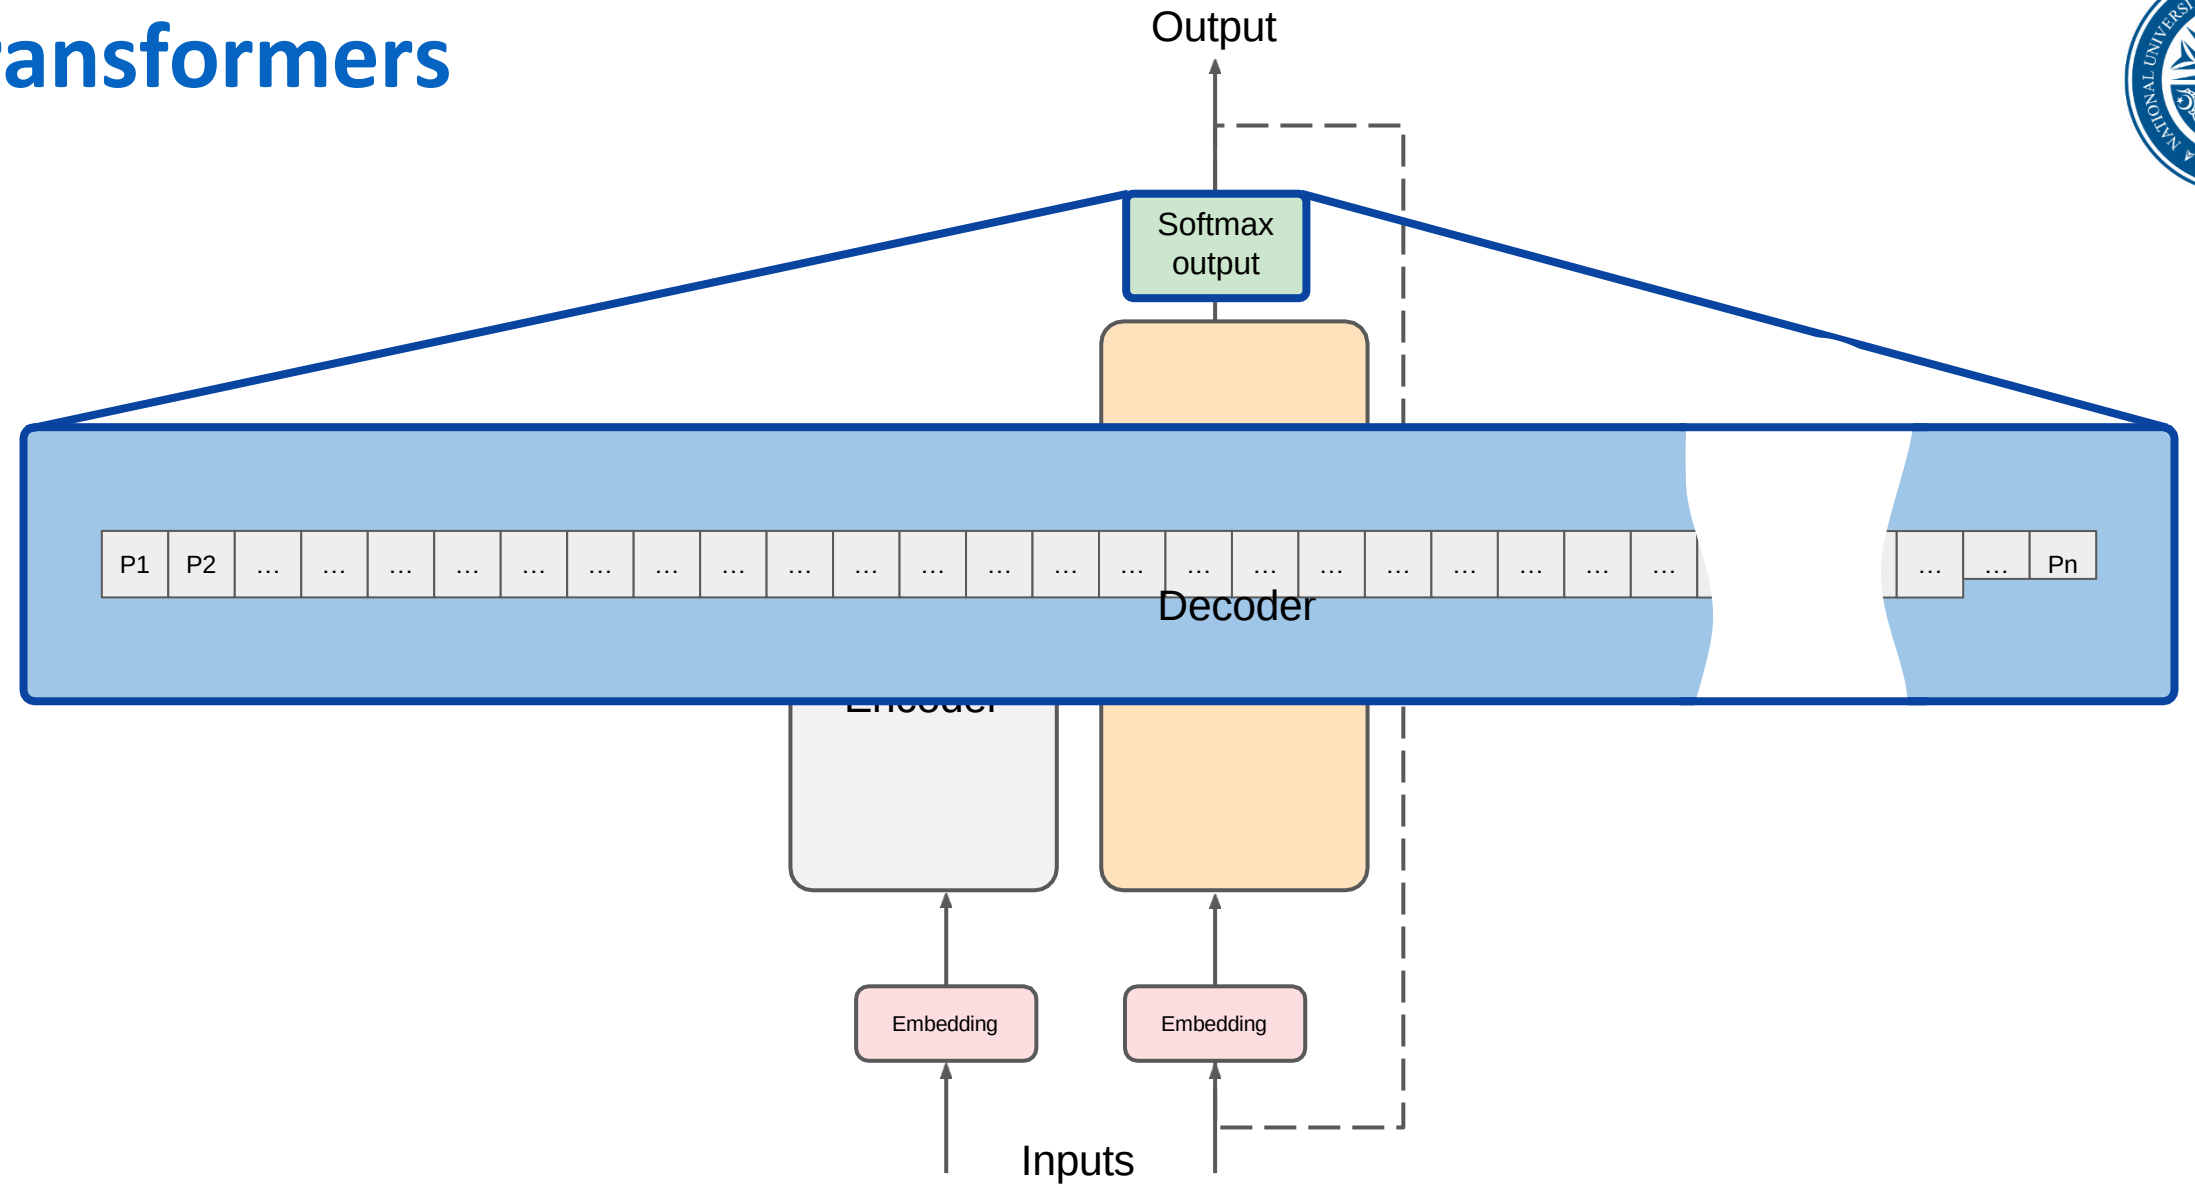

Transformers

Language Generation

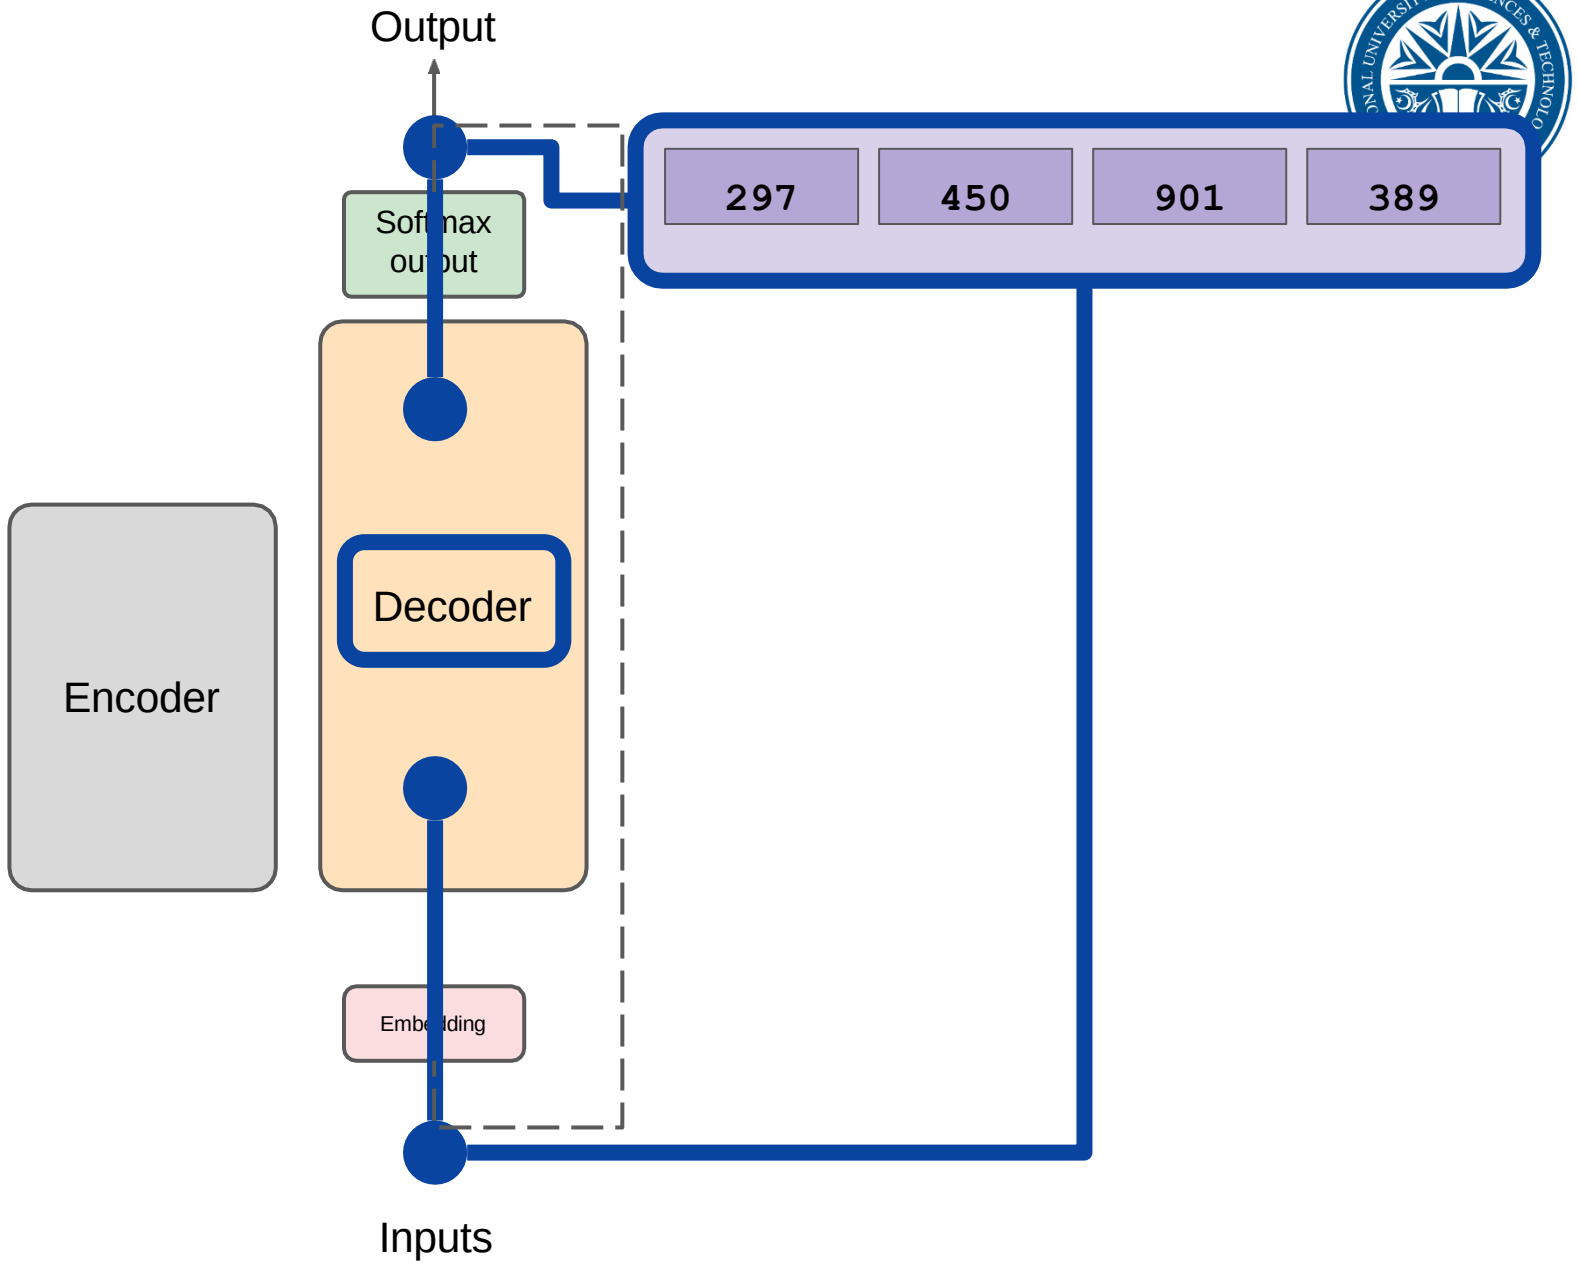

Next word (/token) prediction

Text sample: LLMs learn to predict one word at a time

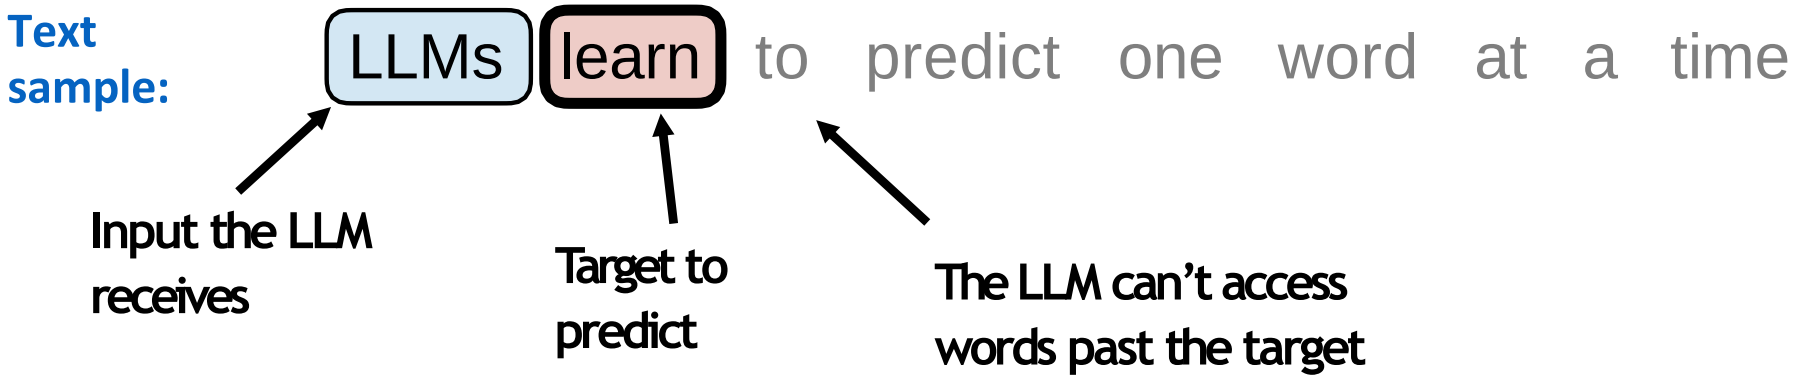
Text sample:
LLMs learn to predict one word at a time
Input the LLM receives
Target to predict
The LLM can't access words past the target

Sample 1 LLMs learn to predict one word at a time

Sample 2 LLMs learn to predict one word at a time

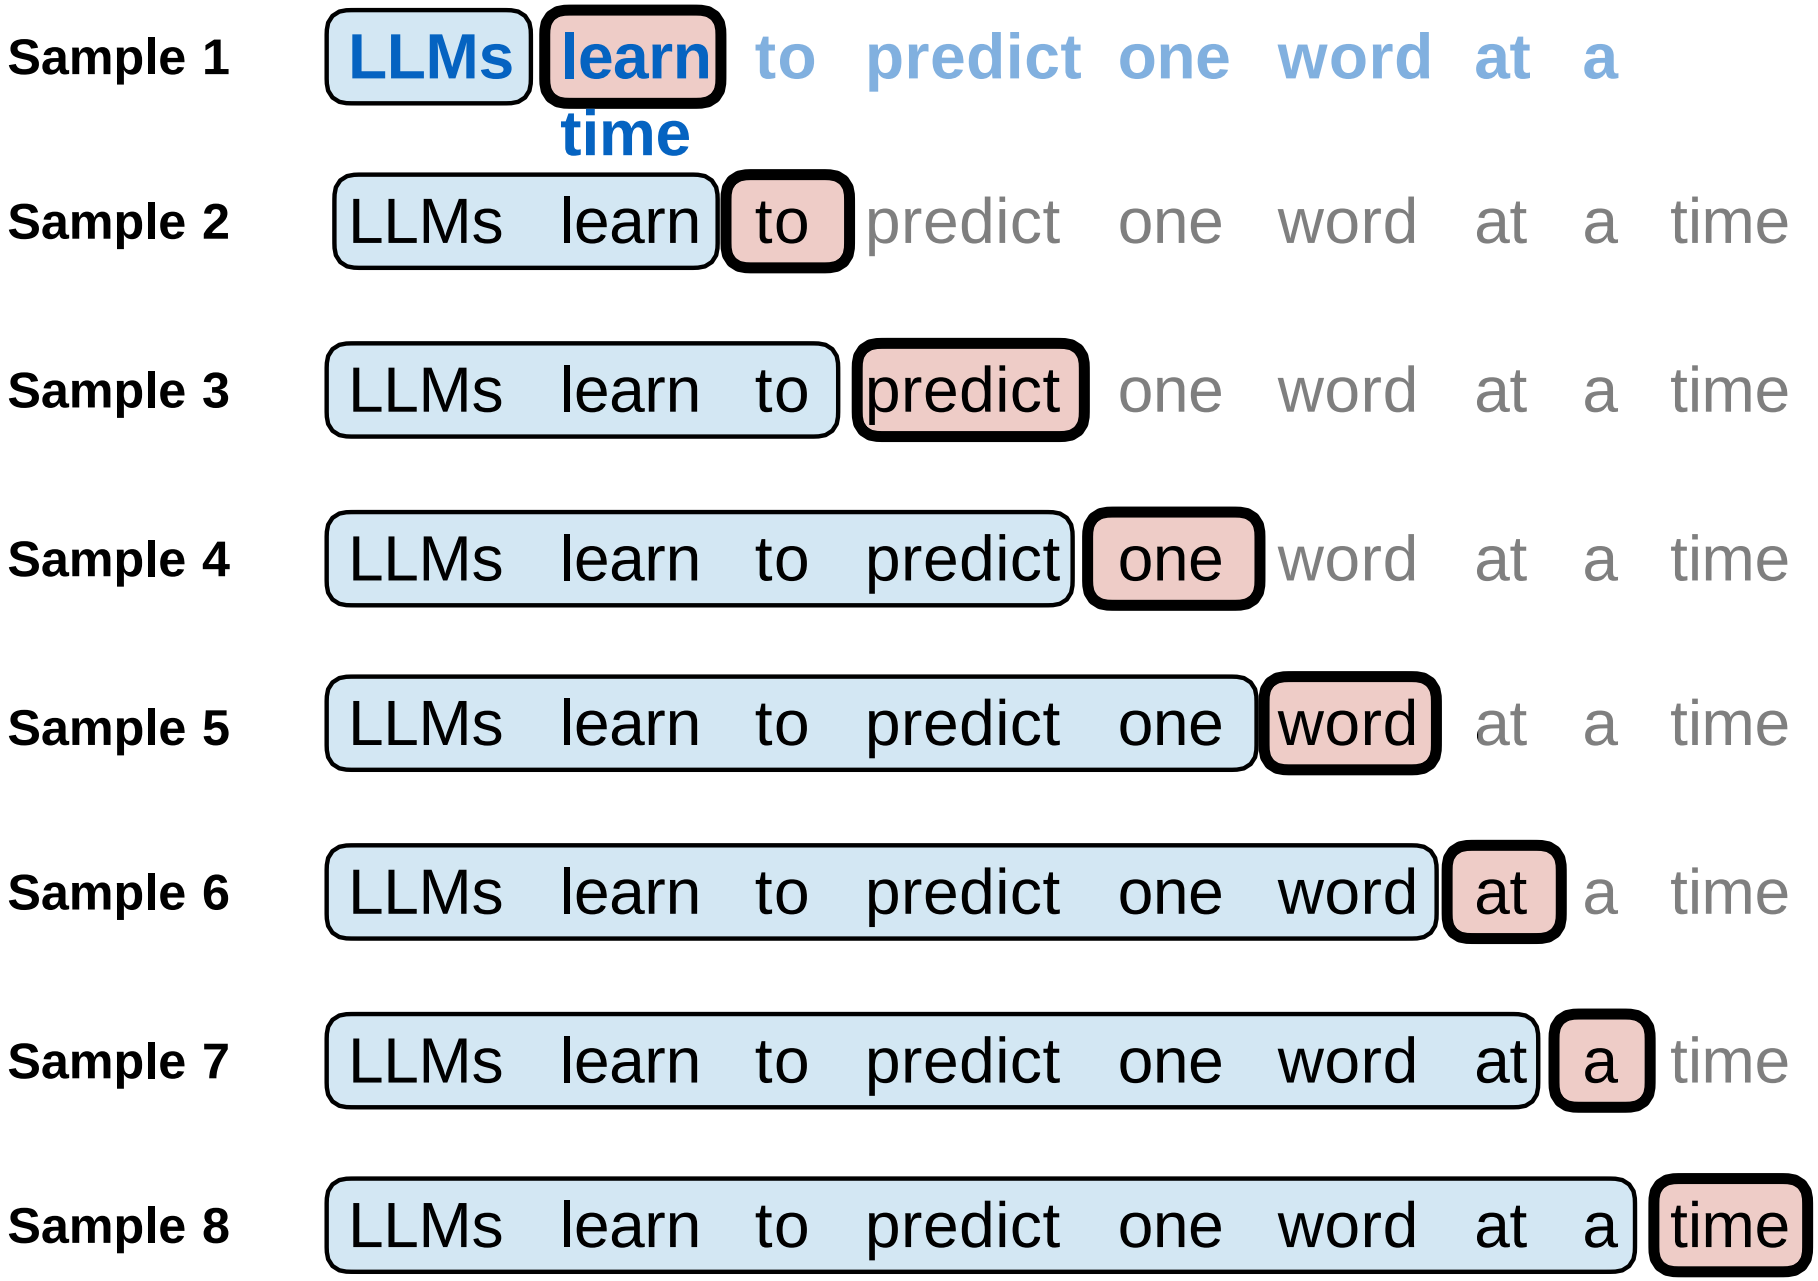
Sample 1
LLMs learn to predict one word at a time
Sample 2
LLMs learn to predict one word at a time
Sample 3
LLMs learn to predict one word at a time
Sample 4
LLMs learn to predict one word at a time
Sample 5
LLMs learn to predict one word at a time
Sample 6
LLMs learn to predict one word at a time
Sample 7
LLMs learn to predict one word at a time
Sample 8
LLMs learn to predict one word at a time

How do LLMs generate multi-word outputs?

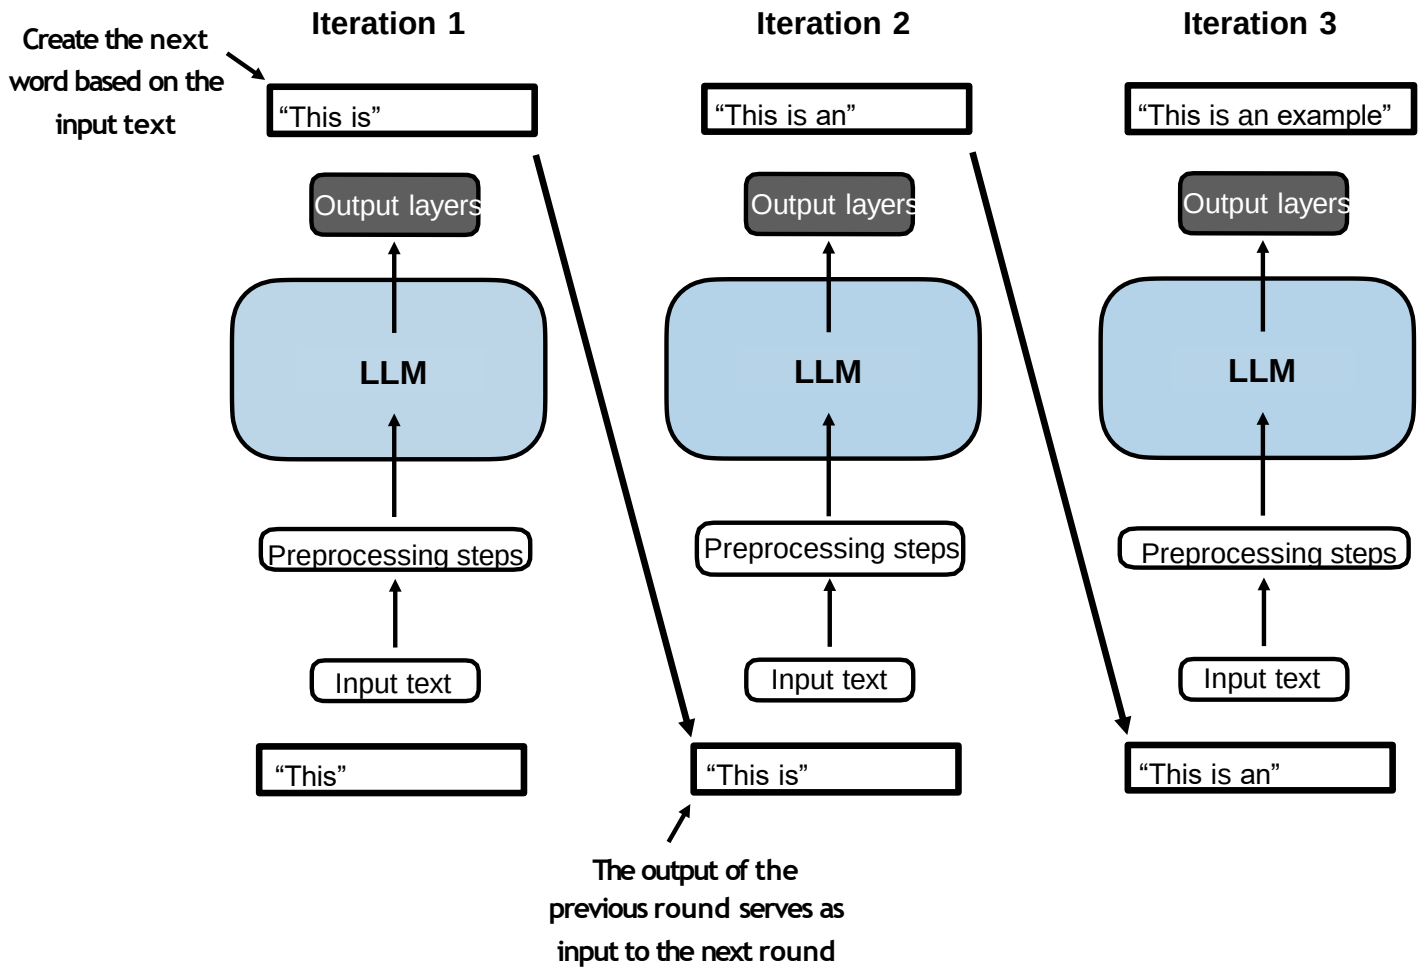

LM Landscape pre GPT-3

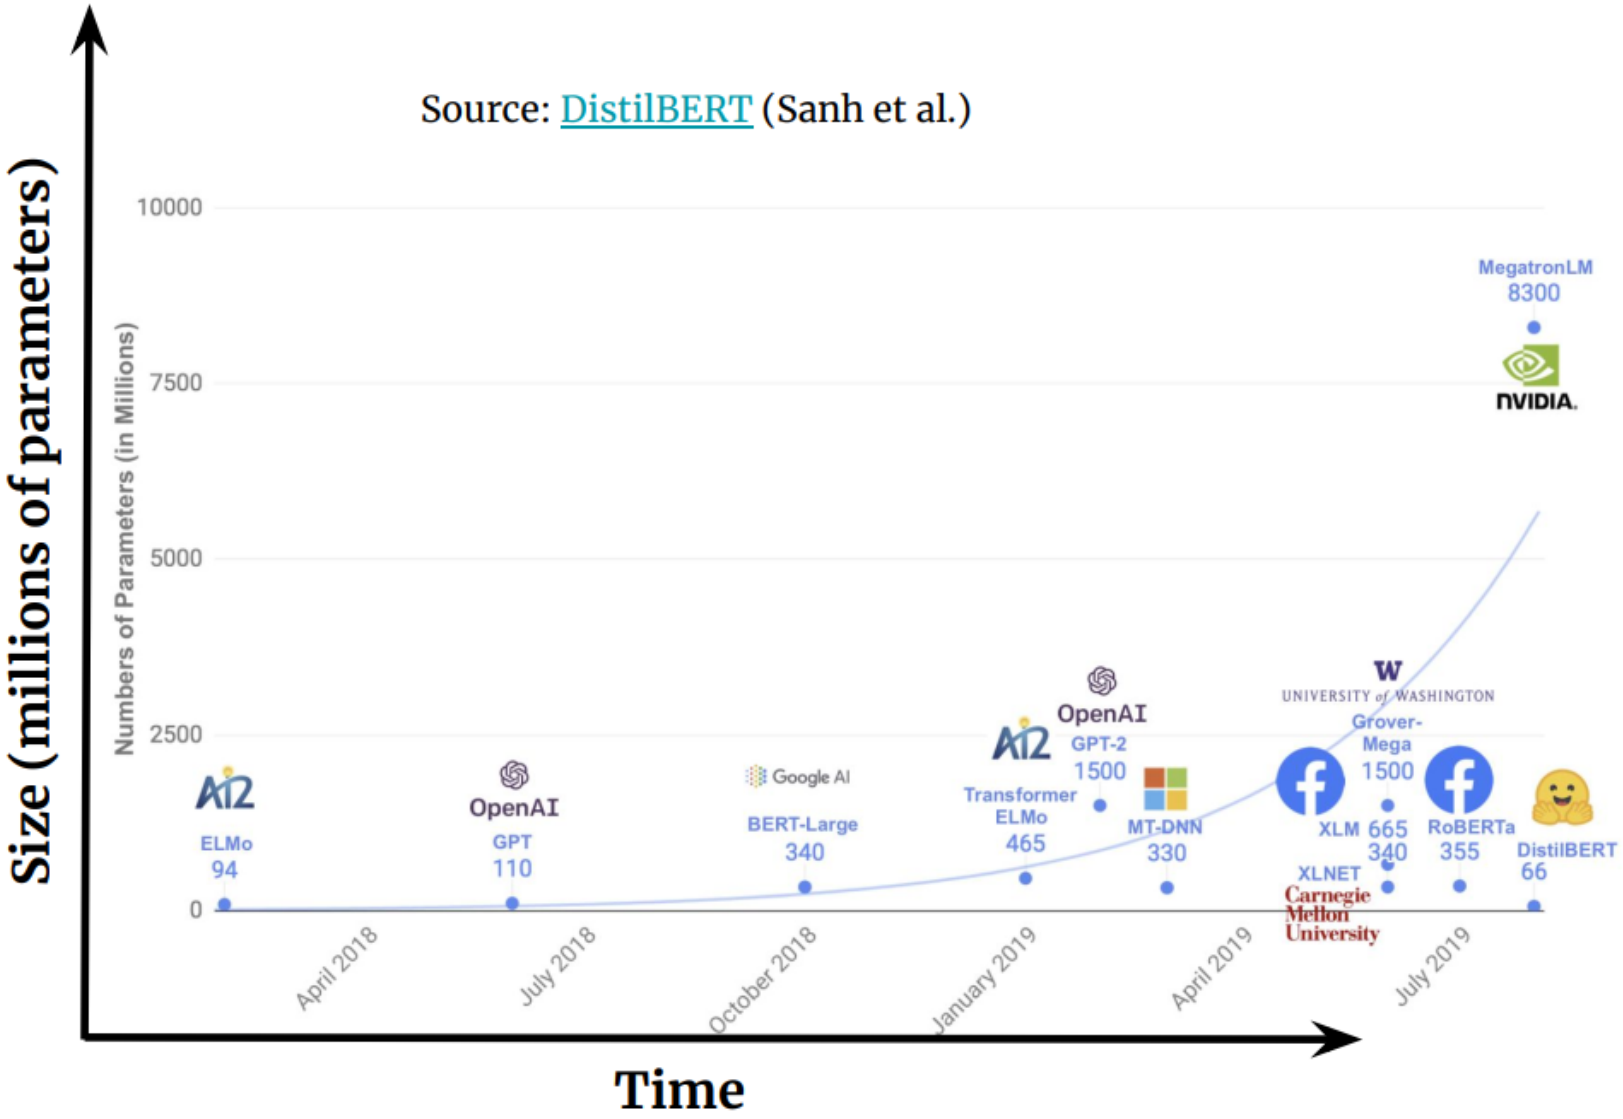

LM Landscape with GPT-3

Source:
https://bmk.sh/2020/05/29/GPT-3-A-Brief-Summary/

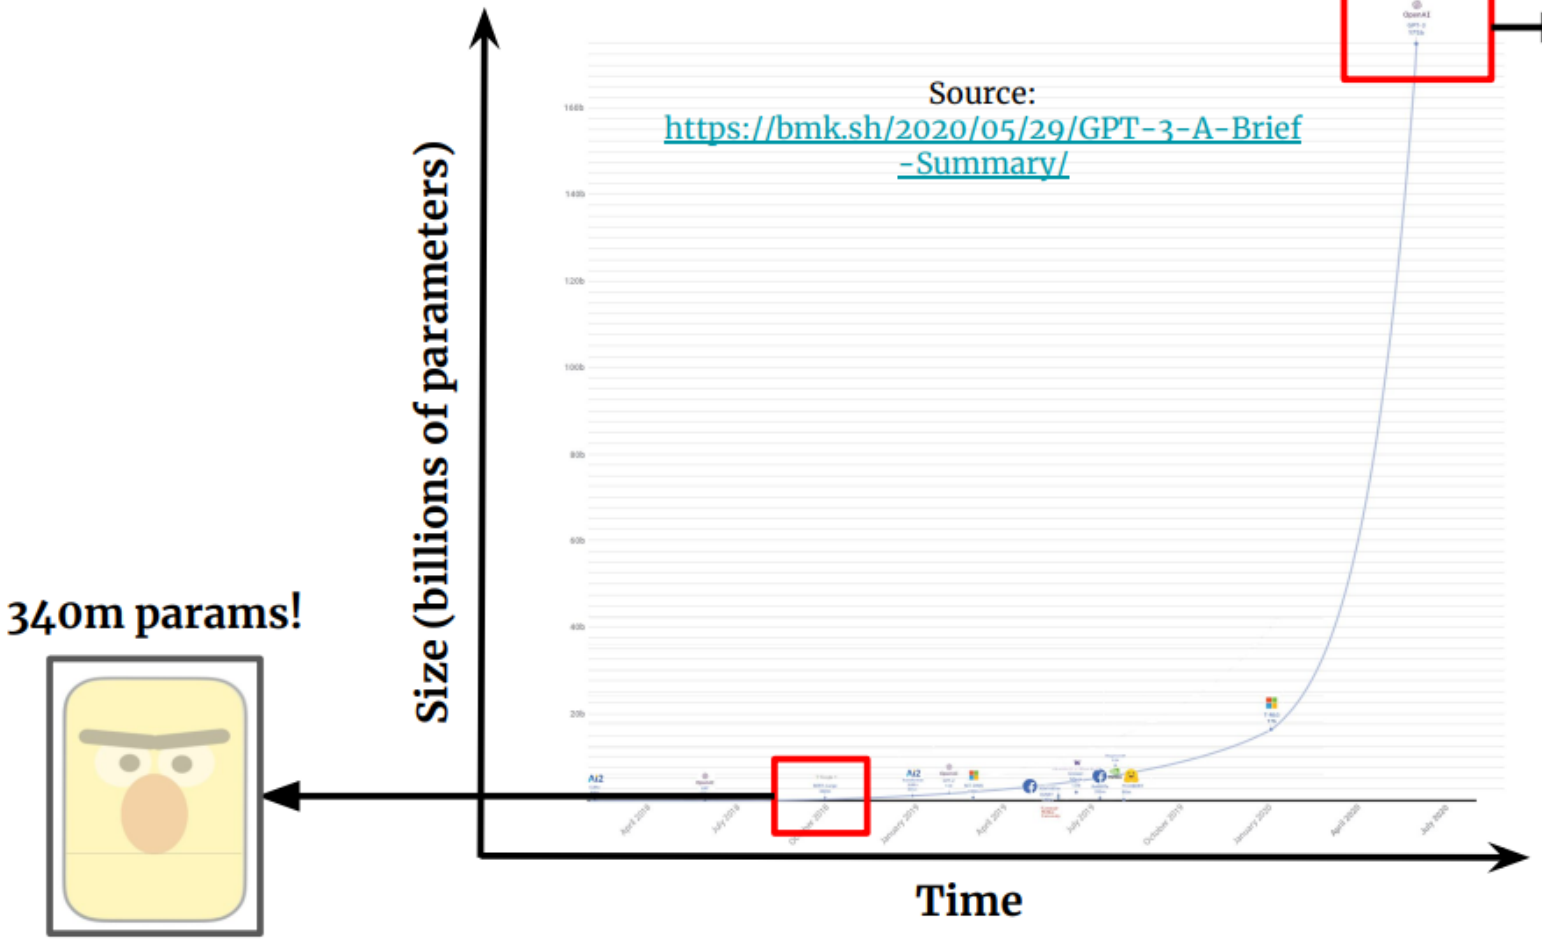

175b params!
GPT-2 was 1.5b

340m params!

Why Scale?

- Study conducted by OpenAI → **Scaling Laws for Neural Language Models** (Kaplan et al. 2020)
- A few **key findings**:
  - Performance depends strongly on scale, weakly on model shape
  - Smooth power laws ($y = ax^k$) b/w empirical performance & N – parameters, D – dataset size, C – compute
  - Transfer improves with test performance
  - Larger models are more sample efficient

Bigger is Better!

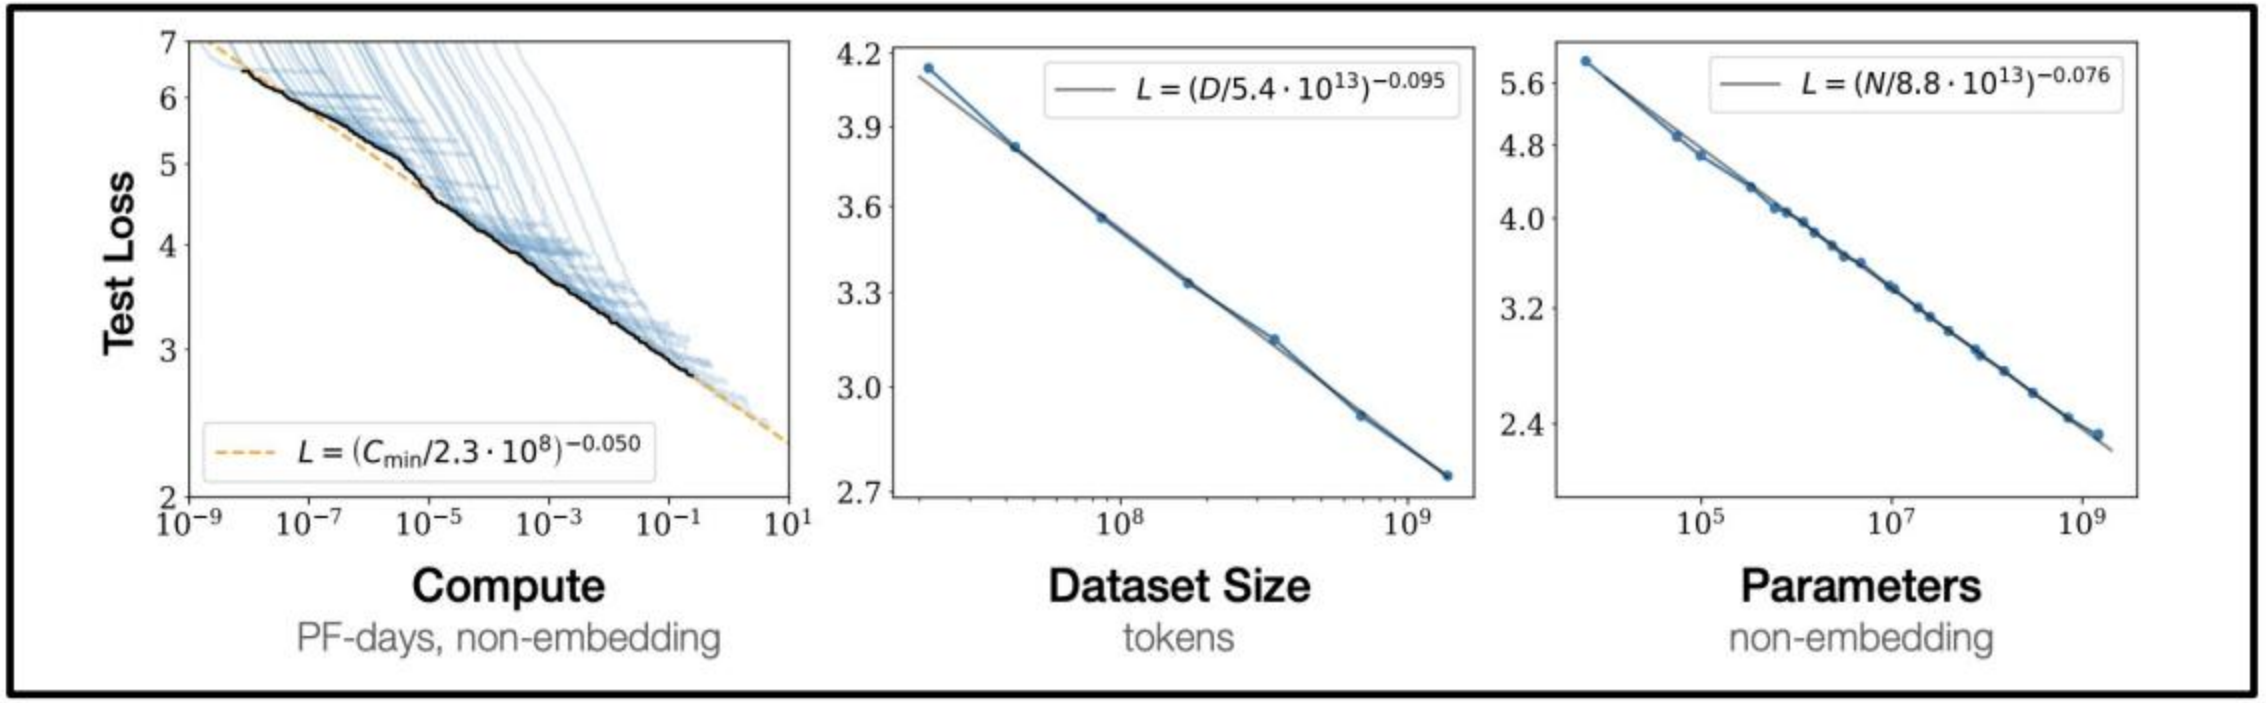

Bigger is Better!

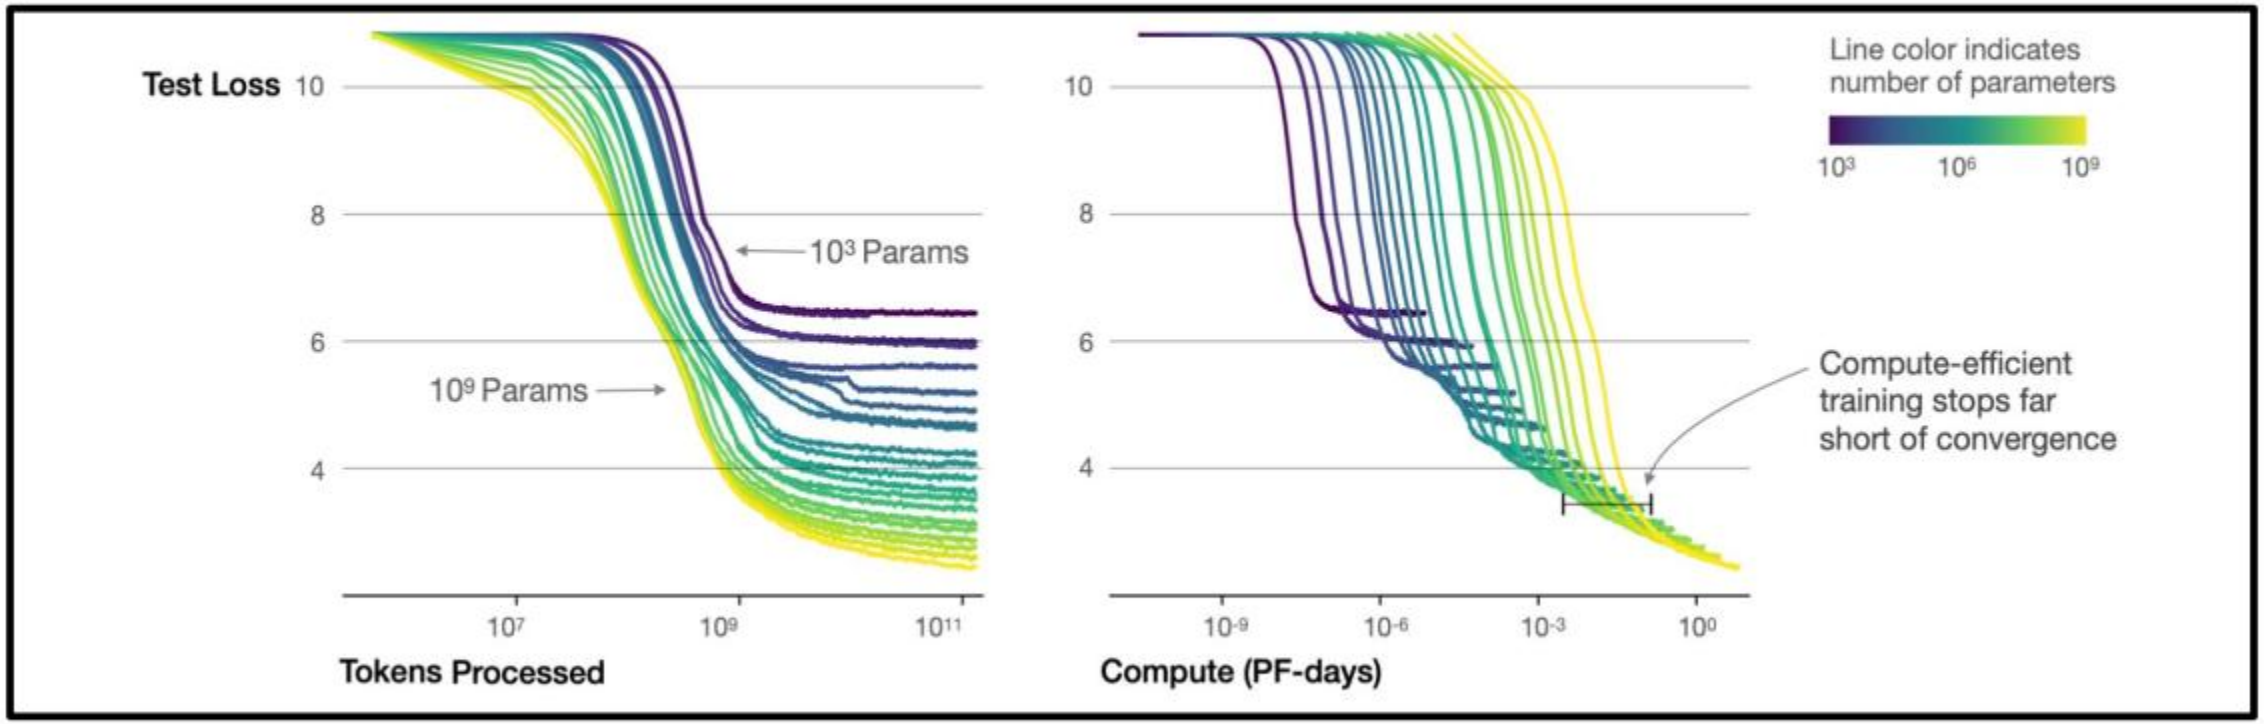

GPT-3 → GPT-2

GPT-3 == A very big GPT-2

- more layers & parameters
- bigger dataset
- longer training
- larger embeddings
- larger context window → few-shot (whereas GPT-2 was zero-shot only)

GPT-3 is MASSIVE!

- 96 decoder blocks (2x GPT-2)
- Context size: 2048 (2x GPT-2)
- Embedding size: 12288 (~8x GPT-2)
- Params: 175b (~117x GPT-2)

GPT-3 is MASSIVE!

Model Name	n_{params}	n_{layers}	d_{model}	n_{heads}	d_{head}	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	6.0×10^{-4}
GPT-3 Medium	350M	24	1024	16	64	0.5M	3.0×10^{-4}
GPT-3 Large	760M	24	1536	16	96	0.5M	2.5×10^{-4}
GPT-3 XL	1.3B	24	2048	24	128	1M	2.0×10^{-4}
GPT-3 2.7B	2.7B	32	2560	32	80	1M	1.6×10^{-4}
GPT-3 6.7B	6.7B	32	4096	32	128	2M	1.2×10^{-4}
GPT-3 13B	13.0B	40	5140	40	128	2M	1.0×10^{-4}
GPT-3 175B or "GPT-3"	175.0B	96	12288	96	128	3.2M	0.6×10^{-4}

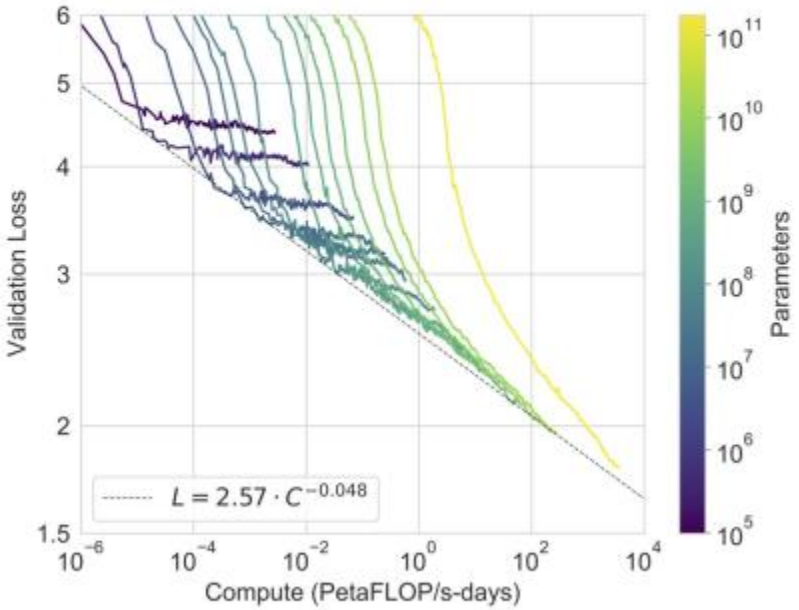

- All models were trained on 300B tokens
- Follows power law argued in Kaplan et al.
- "GPT-3" → GPT-3 175B

In-Context Learning

	No Prompt	Prompt
Zero-shot (0s)	skicts = sticks	Please unscramble the letters into a word, and write that word: skicts = sticks
1-shot (1s)	chiar = chair skicts = sticks	Please unscramble the letters into a word, and write that word: chiar = chair skicts = sticks
Few-shot (FS)	chiar = chair [...] pciinc = picnic skicts = sticks	Please unscramble the letters into a word, and write that word: chiar = chair [...] pciinc = picnic skicts = sticks

GPT training pipeline

Stage	Pretraining	Supervised Finetuning	Reward Modeling	Reinforcement Learning
Dataset	**Raw internet** text trillions of words low-quality, large quantity	**Demonstrations** Ideal Assistant responses, ~10-100K (prompt, response) written by contractors low quantity, high quality	**Comparisons** 100K –1M comparisons written by contractors low quantity, high quality	**Prompts** ~10K-100K prompts written by contractors low quantity, high quality
Algorithm	**Language modeling** predict the next token	**Language modeling** predict the next token	**Binary classification** predict rewards consistent w preferences	**Reinforcement Learning** generate tokens that maximize the reward
		init from	init from	init from SFT use RM
Model	**Base model**	**SFT model**	**RM model**	**RL model**
Notes	1000s of GPUs months of training ex: GPT, LLaMA, PaLM can deploy this model	1-100 GPUs days of training ex: Vicuna-13B can deploy this model	1-100 GPUs days of training	1-100 GPUs days of training ex: ChatGPT, Claude can deploy this model

Credits : @karpathy

Generative configuration parameters for inference

Generative configuration - inference parameters

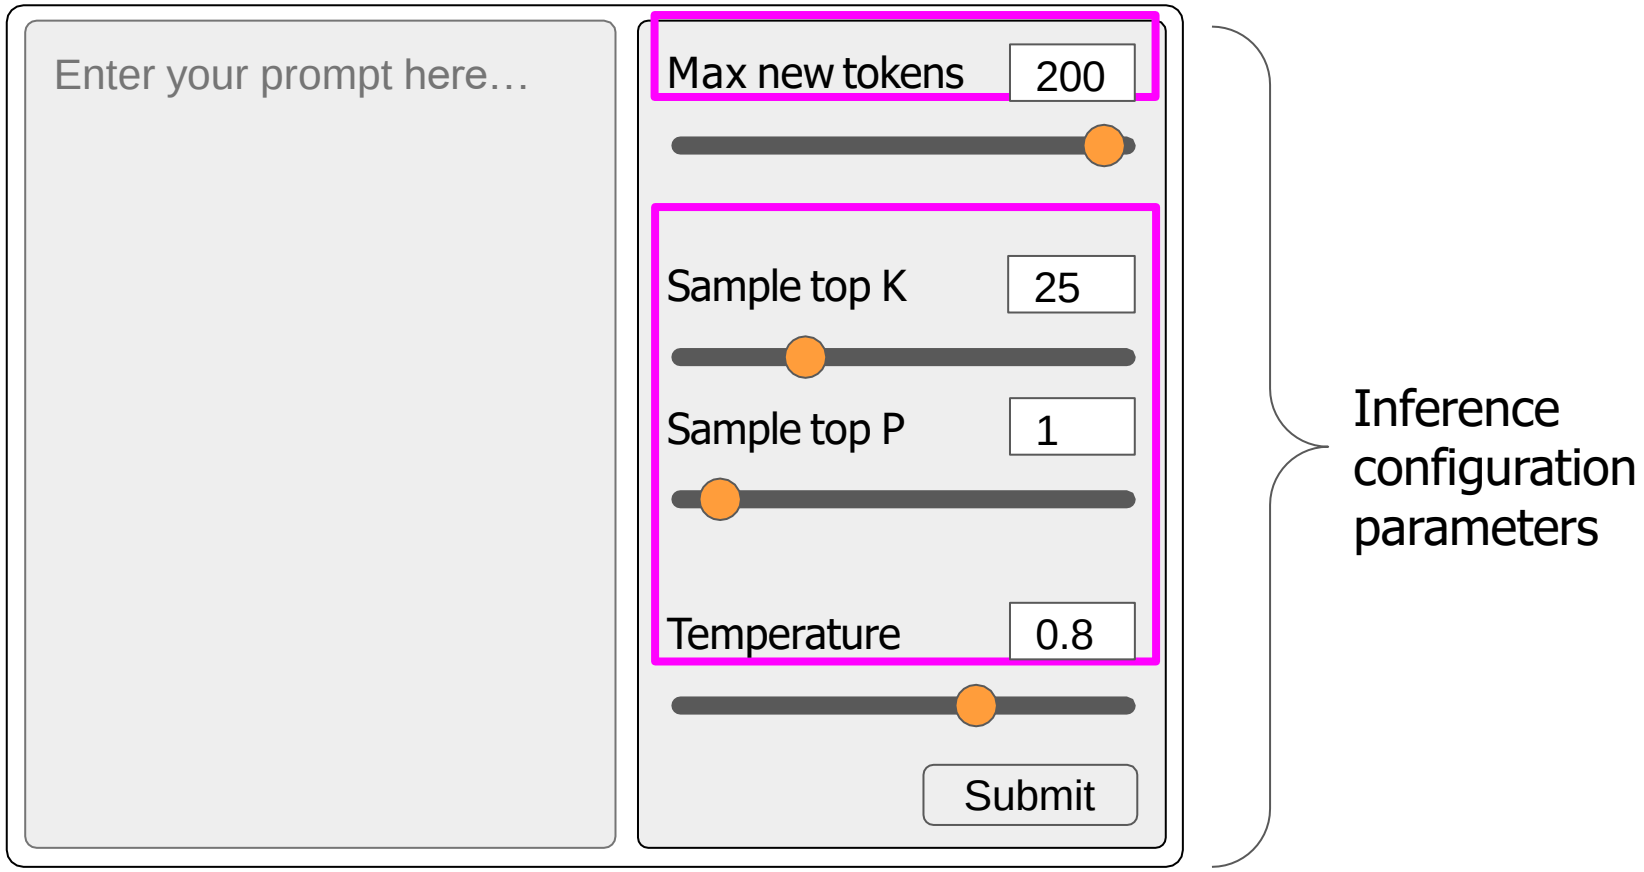

Fine-tuning an LLM with instruction prompts

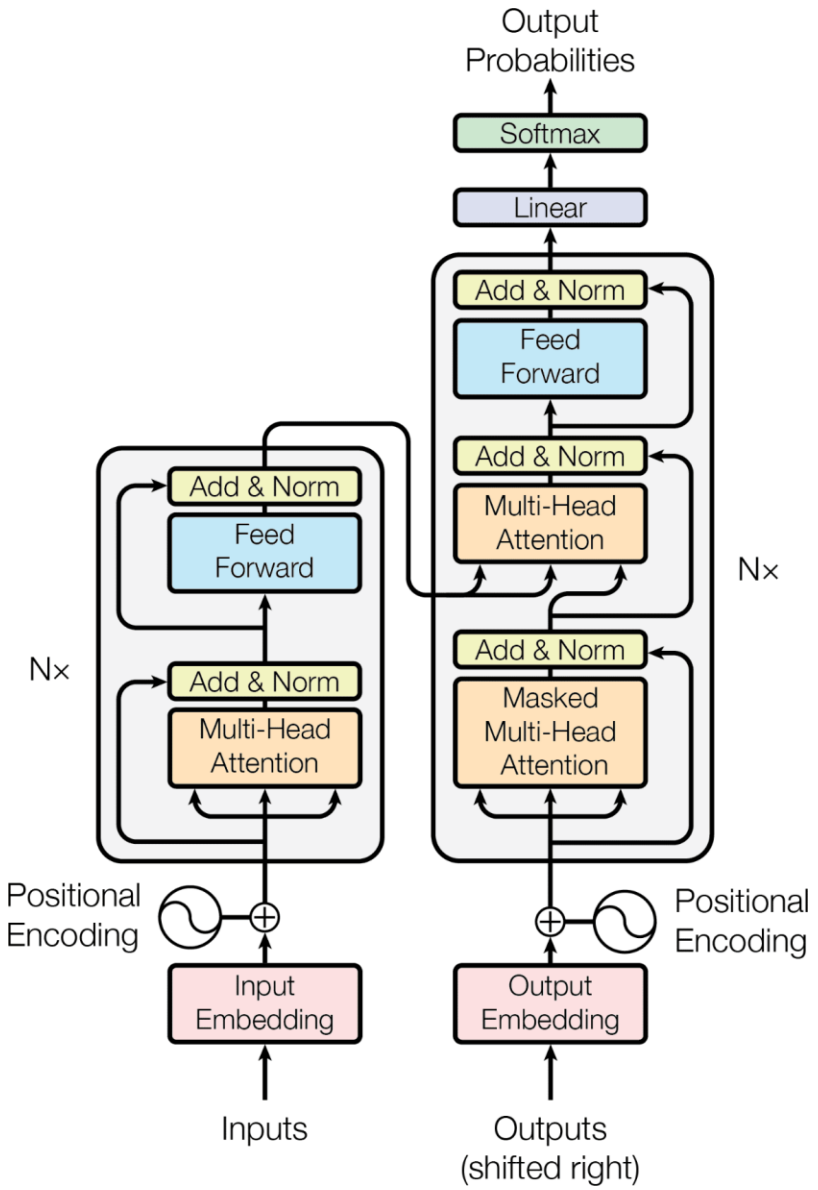

In-context learning (ICL) - zero shot inference

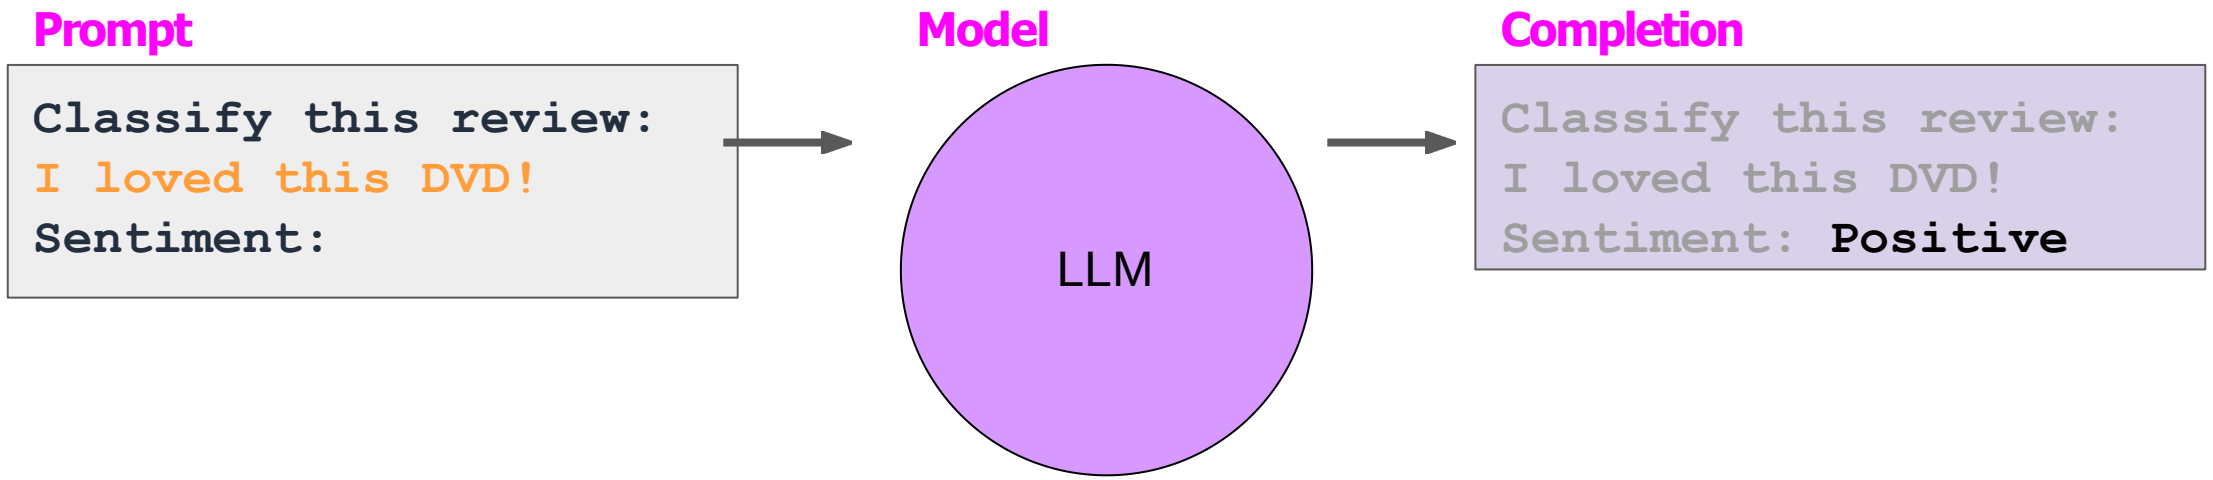

LLM fine-tuning at a high level

LLM pre-training

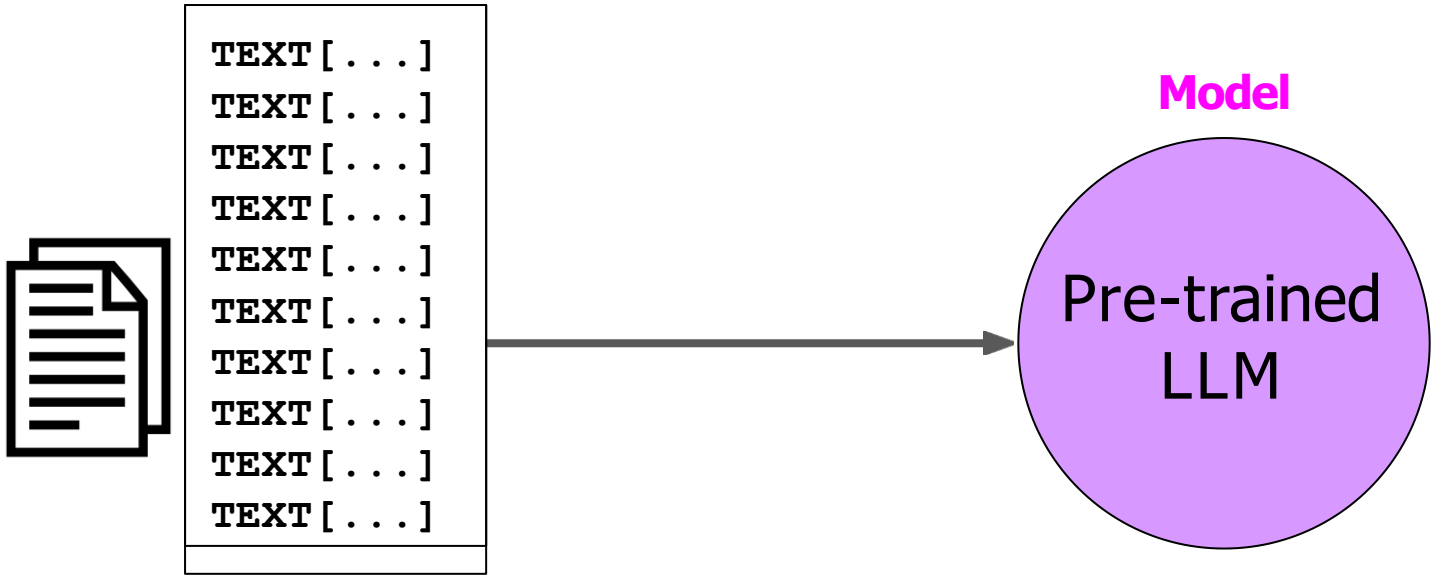

GB - TB - PB
of unstructured textual data

LLM fine-tuning at a high level

LLM fine-tuning

GB - TB
of labeled examples for a specific task or set of tasks

LLM fine-tuning at a high level

LLM fine-tuning

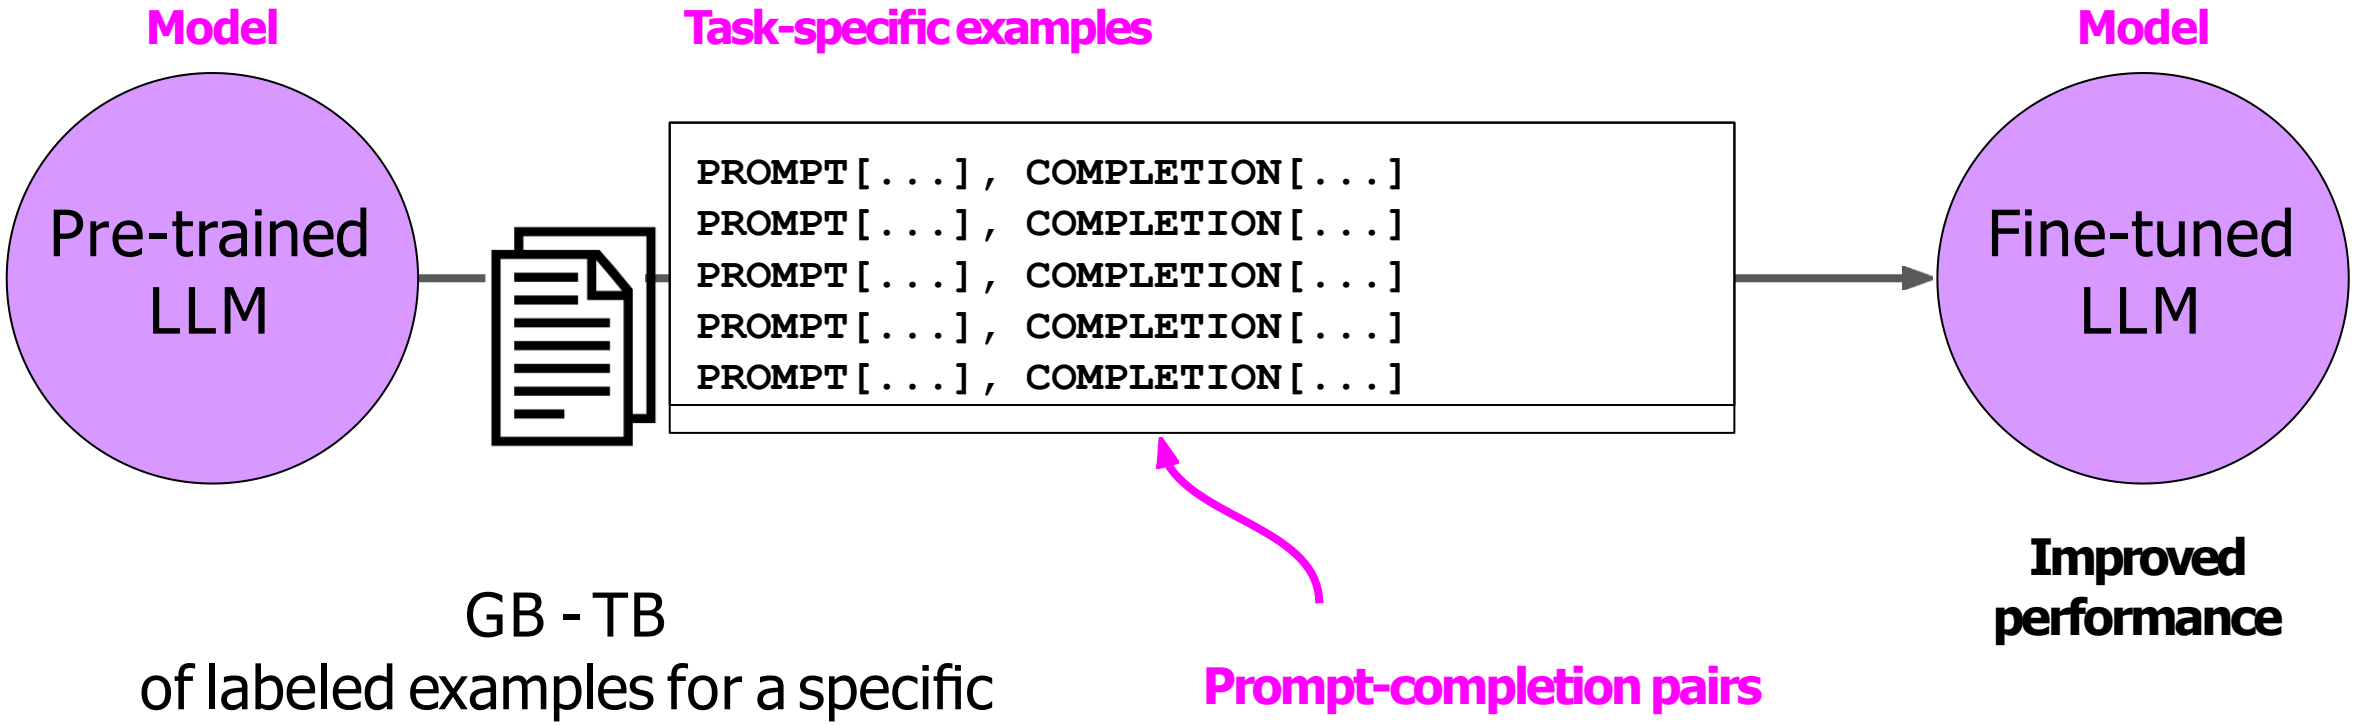

Multi-task, instruction fine-tuning

Multi-task, instruction fine-tuning

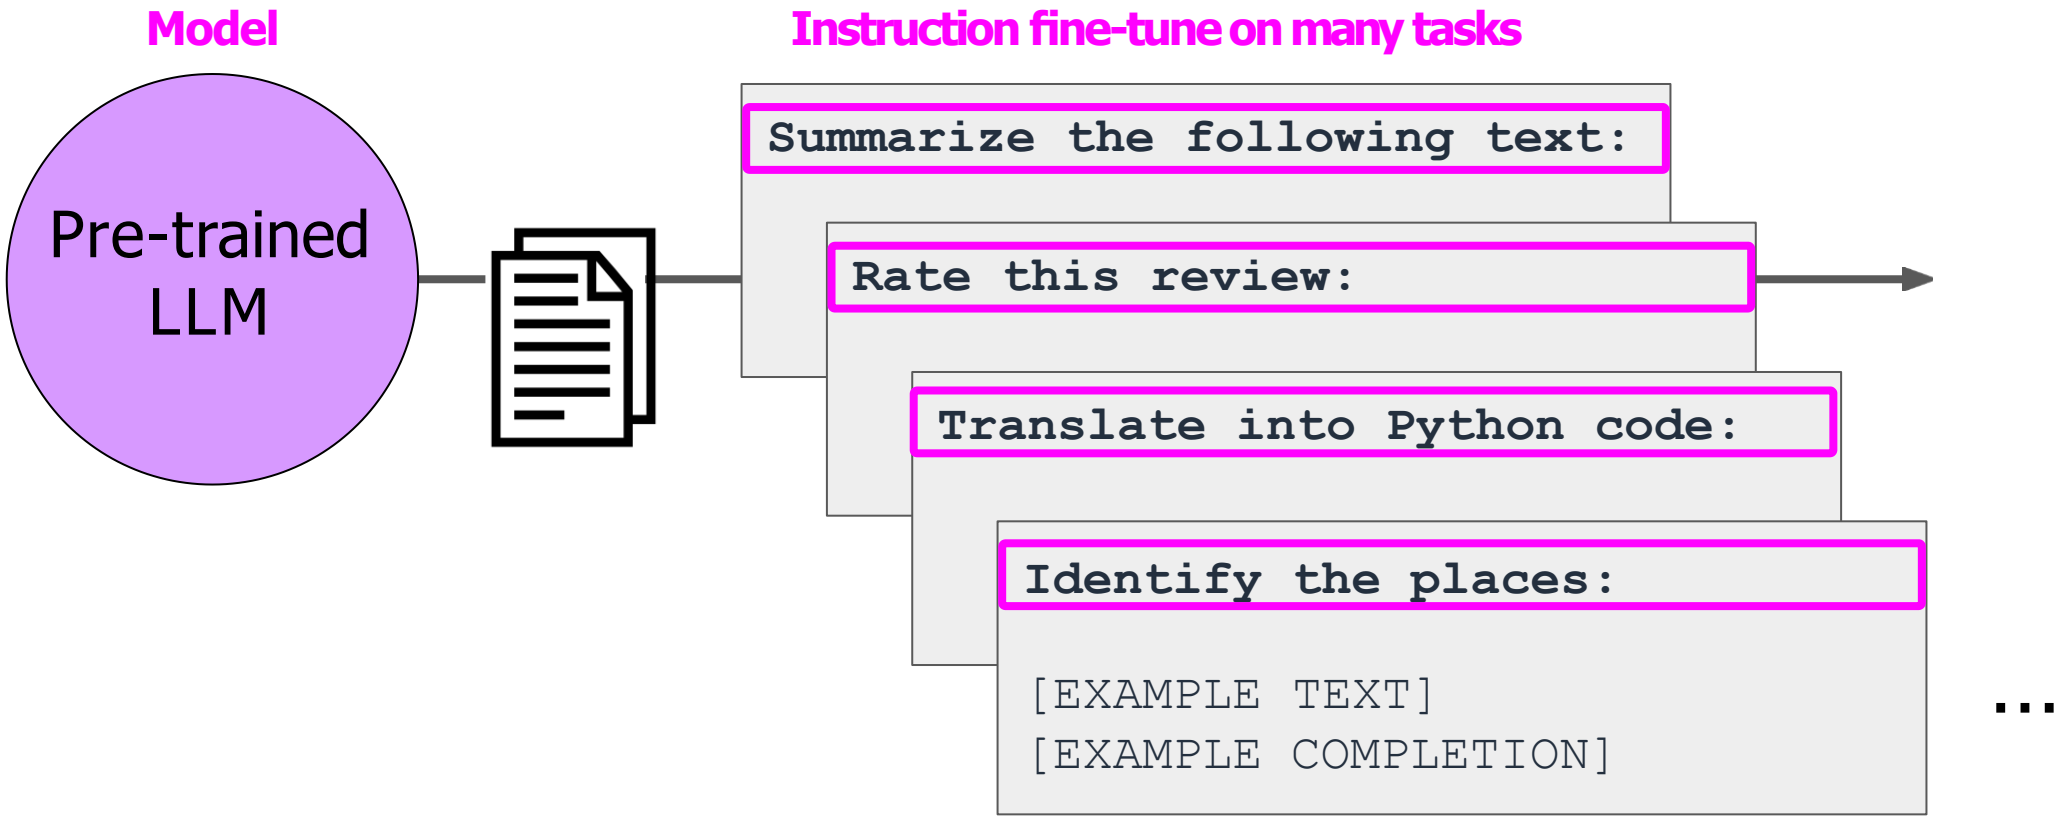

Multi-task, instruction fine-tuning